

Lab #3. Cache Lab

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering

Sogang University

Change of Plan: Labs

- Labs will count for 30% of the total score in semester
- We have done three lab assignments so far
 - Lab #1: Bit Lab (6%)
 - Lab #2: Reversing Lab (12%)
 - Lab #3: Cache Lab (12% \Rightarrow 8%)
 - I fixed the problem to make it easier than my original plan
 - And its weight in the total score has been slightly reduced
 - Remaining 4% point will be given as free bonus point

About this lab

- In this lab, you have to **write a C program that simulates cache memory access**
 - The cache parameters (e.g., # of lines) will be given as input
 - For the provided memory access sequence, your program must decide whether each access will be a **cache hit** or **cache miss**
- To do this assignment, you must clearly understand the **operation of cache memory**
 - Also, it is a good chance to practice implementing certain idea into concrete code

Remind: Cheating Policy

- **Cheating in assignment will give you a serious penalty**
 - Your final grade will be downgraded (e.g., from **B+** to **C+**)
- **Scope of cheating in assignment**
 - Copying the code of other people
 - Sharing your solution with others
 - Asking ChatGPT to write your code
 - Discussing with others how to solve the problem

General Information

■ Check the *Assignment* tab of *Cyber Campus*

- Skeleton code (`Lab3.tgz`) is attached together with this slide
- Submission will be accepted in the same post, too

■ Deadline: 6/26 Wednesday 23:59

- **No late submission for this Lab!**

■ Please read the instructions in this slide carefully

- This slide is a step-by-step tutorial for the lab
- It also contains important submission guidelines
 - If you do not follow the guidelines, **you will get penalty**

Skeleton Code Structure

- Copy Lab3.tgz into CSPRO server and decompress it
 - Recommend to use cspro2.sogang.ac.kr (Ubuntu 20.04)
 - Don't decompress-and-copy; copy-and-decompress
- **cache-simulator**: In Lab #3, this is the only problem directory that you have to care about
- **check.py**: Script for self-grading (explained later)
- **config**: Used by grading script (you may ignore)

```
jason@ubuntu:~$ tar -xzf Lab3.tgz
jason@ubuntu:~$ ls Lab3
cache-simulator  check.py  config
```

Problem Directory

- **cache.h**: Provides type definitions and function declarations for the cache simulator
- **cache.c**: The actual cache simulator logic that you have to fill in (do **NOT** modify any other file)
- **main.c**: Program to run and test your simulator code
- **Makefile**: Already given for you (just type make to build)
- **testcase**: Contains test cases and expected outputs

```
jason@ubuntu:~/Lab3$ cd cache-simulator/  
jason@ubuntu:~/Lab3/cache-simulator$ ls  
cache.c  cache.h  main.c  Makefile  testcase
```

Input of Cache Simulator

- Let's first check the format of inputs and outputs
- **config-N**, **trace-N** are the inputs of the cache simulator
 - **config-*** contains the cache parameters **b**, **E**, and **s** (be careful on the **lower/upper case**)
 - Assume that $0 < \mathbf{b} \leq 16$, $0 < \mathbf{E}$, $0 < \mathbf{s} \leq 16$
 - **trace-*** contains a sequence of memory addresses to access
 - Assume that all accesses are **1-byte access**

```
jason@ubuntu:~/Lab3/cache-simulator/testcase$ cat config-1
2 1 4
jason@ubuntu:~/Lab3/cache-simulator/testcase$ cat trace-1
BA00
BA04
AA08
BA05
```


Output of Cache Simulator

- Your compiled program must decide whether each memory access will result in cache hit or cache miss
 - Ex) The simulator below is saying that the first three accesses (BA00, BA04, AA08) are **miss** and the next one (BA05) is **hit**
- **ans-N** is the expected output for **config-N** & **trace-N**

```
jason@ubuntu:~/Lab3/cache-simulator$ make
gcc -fsanitize=address cache.c main.c -o main.bin
jason@ubuntu:~/Lab3/cache-simulator$ cd testcase/
jason@ubuntu:~/Lab3/cache-simulator/testcase$ ../main.bin ./config-1 trace-1
MISS
MISS
MISS
HIT
```

Must match with the
content of **ans-1** file

Driver Code (main.c)

- The **main.c** file in skeleton code is already doing many things for the cache simulator
 - It reads in the cache parameters (**b**, **E**, and **s**) from **config-N** and initialize the cache structure by calling **init_cache()**
 - Then, it iteratively reads in each line of **trace-N** and calls **access_memory()** to decide whether it's cache hit or miss
 - Lastly, it calls **free_cache()** to clean up the heap memory
- You must first take enough time to carefully read and understand the code in **main.c**

Tasks to do

- You have to implement the three functions below
 - Read the type definitions and comments in **cache.h** carefully
- Depending on the value of cache parameters, your cache can be direct-mapped or set/fully associative
 - And your cache must implement **LRU policy** for eviction

```
// Allocate and initialize cache_t structure. (...)
cache_t* init_cache(int b, int E, int s);

// Simulate memory access on address 'addr'. Return 1 if
// the access is a cache hit and 0 if it's a miss. (...)
int access_memory(cache_t* cache, addr_t addr);

// Free the dynamically allocated memory in 'cache'. (...)
void free_cache(cache_t* cache);
```

What is this error message?

- You may encounter this kind of error message when you run the compiled **main.bin** program
- In our **Makefile**, I used some special compiler flag when building the **main.bin** binary
 - As shown below, it will report buffer overflows or memory leaks (unallocated memory when **main()** returns) during the runtime

```
jason@ubuntu:~/Lab3/cache-simulator/testcase$ ../main.bin ./config-1 trace-1
MISS
MISS
MISS
HIT
=====
==2289==ERROR: LeakSanitizer: detected memory leaks
```

Self-Grading

- Once you think everything is done, run **check.py** to confirm that you pass all the provided test cases
 - Each character in the result has following meaning:
 - 'O': correct, 'X': wrong, 'C': compile error,
 - 'T': timeout, 'E': runtime error
- You'll get the point based on the number of test cases that your code passes (100 points in total)

```
jason@ubuntu:~/Lab3$ ls
cache-simulator  check.py  config
jason@ubuntu:~/Lab3$ ./check.py
[*] cache-simulator : 0000
```

Submission Guideline

- You should submit only one C source file
 - `cache.c` (be careful not to submit `cache.h`)
- If the submitted file does not compile by typing "make" command, **cannot give you any point** for that problem
- Submission format
 - Upload this file directly to *Cyber Campus* (**do not zip it**)
 - **Do not change the file name** (e.g., adding any prefix or suffix)
 - If your submission format is wrong, you will get **-20% penalty**